

REVISITING SPARSE MATRIX COLORING AND BICOLORING

Alexis Montoison, Guillaume Dalle



November, 17th 2025
JuMP-dev
Auckland

- Automatic Differentiation (AD) is at the core of modern scientific computing and nonlinear optimization.
- Solvers such as Ipopt, Knitro, Uno, and MadNLP require Jacobians and Lagrangian Hessians at every iteration in order to compute search directions.
- These derivative matrices are large but sparse, and exploiting this sparsity is key to efficient AD and linear algebra.
- Coloring and bicoloring provides an elegant way to reduce the number of AD passes needed to recover sparse Jacobians and Hessians.

Jacobian computation via automatic differentiation

- Consider $c : \mathbb{R}^n \rightarrow \mathbb{R}^m$ with Jacobian $J_c(x) = \partial c(x)$.
- Forward-mode AD computes Jacobian-vector products: $u \mapsto J_c(x)u$.
- Reverse-mode AD computes vector-Jacobian products: $v \mapsto v^\top J_c(x)$.
- The full Jacobian can be reconstructed column-wise or row-wise.

Exploiting sparsity

- If columns have disjoint non-zeros, they can be recovered together.

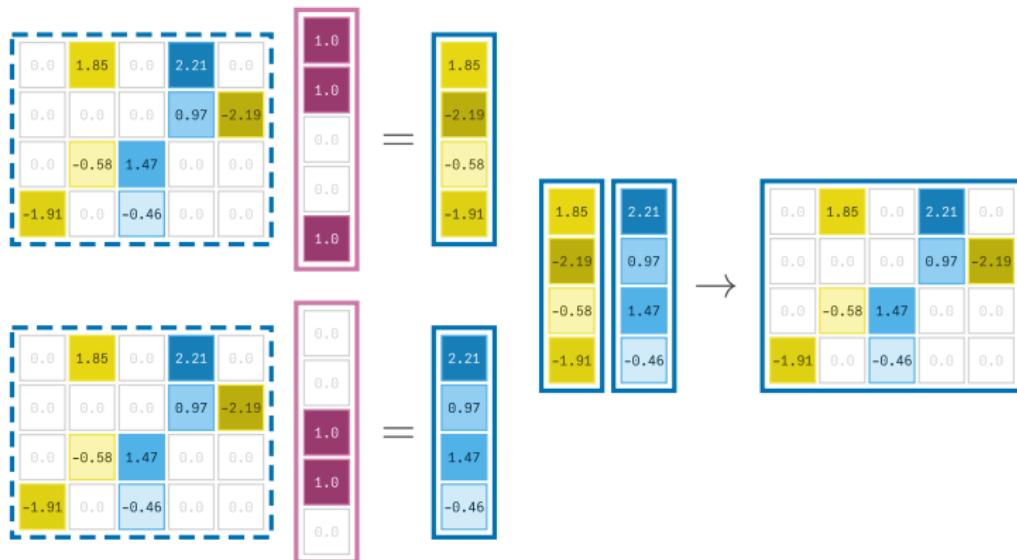


Figure – Materializing a Jacobian with forward-mode AD : (left) compressed evaluation of orthogonal columns (right) decompression to Jacobian matrix.

Exploiting sparsity

- If rows have disjoint non-zeros, they can be recovered together.

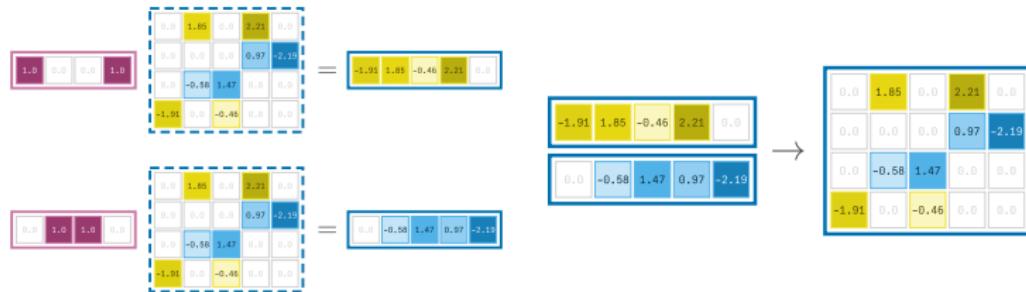


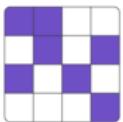
Figure – Materializing a Jacobian with reverse-mode AD : (left) compressed evaluation of orthogonal rows (right) decompression to Jacobian matrix.

Automatic sparse differentiation

(a) AD code transformation

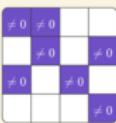
$$f(\mathbf{x}) \rightarrow \text{VJP}(\mathbf{x}, \mathbf{v})$$


(b) Standard AD Jacobian computation

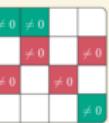
$$\begin{aligned}\text{VJP}(\mathbf{x}, \mathbf{e}_1) &= \begin{array}{c} \text{purple} \\ \text{white} \end{array} \\ \text{VJP}(\mathbf{x}, \mathbf{e}_2) &= \begin{array}{c} \text{white} \\ \text{purple} \end{array} \\ \text{VJP}(\mathbf{x}, \mathbf{e}_3) &= \begin{array}{c} \text{purple} \\ \text{white} \end{array} \\ \text{VJP}(\mathbf{x}, \mathbf{e}_4) &= \begin{array}{c} \text{white} \\ \text{purple} \end{array}\end{aligned}$$


→

(c) ASD Jacobian computation

$$f \rightarrow$$


→



→

$$\begin{aligned}\text{VJP}(\mathbf{x}, \mathbf{e}_1 + \mathbf{e}_4) &= \begin{array}{c} \text{green} \\ \text{white} \\ \text{green} \end{array} \\ \text{VJP}(\mathbf{x}, \mathbf{e}_2 + \mathbf{e}_3) &= \begin{array}{c} \text{red} \\ \text{red} \\ \text{red} \end{array}\end{aligned}$$


→



① Pattern detection

② Coloring

③ Matrix-vector products

④ Decompression

Graph coloring

- This grouping problem can be reformulated as graph coloring.
- The goal is to minimize the number of AD evaluations.
- Graph coloring determines independent sets of columns (or rows).
- Complexity scales with the number of colors instead of n or m .

Graph coloring

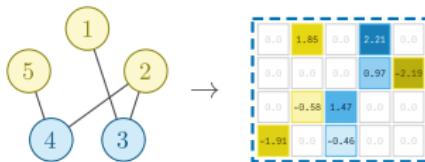


Figure – Optimal graph coloring.

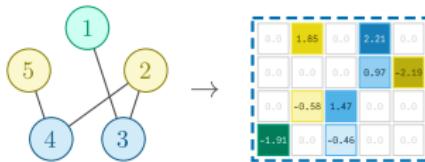


Figure – Suboptimal graph coloring (vertex 1 could be colored in yellow).

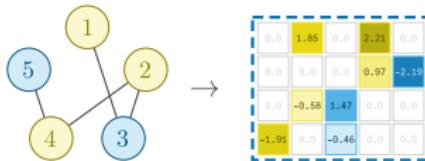


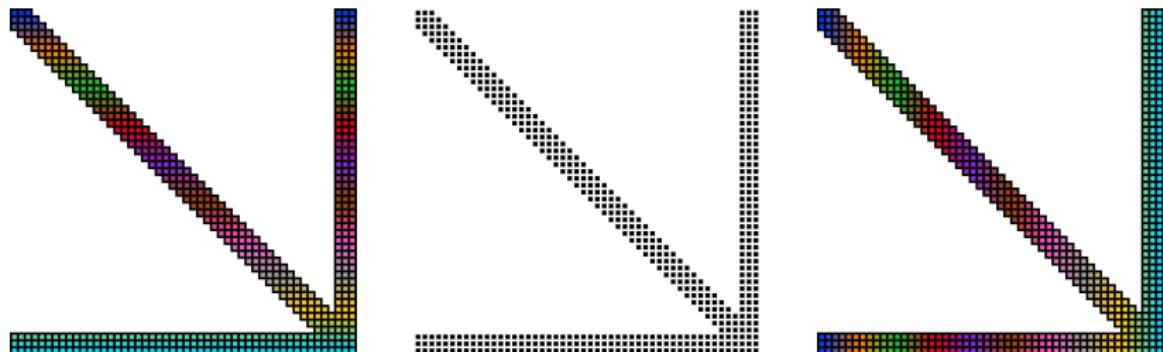
Figure – Infeasible graph coloring (vertices 2 and 4 are adjacent on the graph, but share a color)

Hessian computation via automatic differentiation

- Consider $f : \mathbb{R}^n \rightarrow \mathbb{R}$ with Hessian $H_f(x) = \partial \nabla f(x)$.
- Forward-over-reverse mode AD computes Hessian-vector products: $w \mapsto H_f(x)w$.
- We can exploit symmetry to recover non-zeros from the lower or upper triangle.
- Star coloring and acyclic coloring are the most common symmetric colorings.

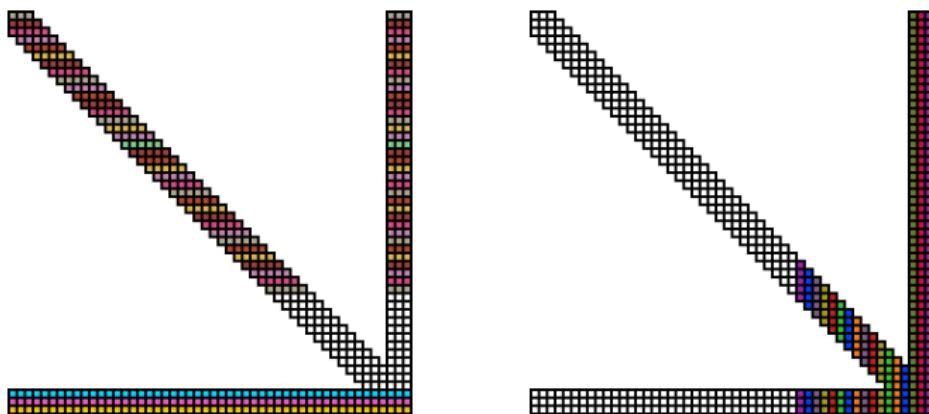
- Some Jacobians contain dense substructures.
- Standard unidirectional coloring can become inefficient.
- Bicoloring jointly colors rows and columns.
- This combines forward- and reverse-mode AD for better performance.

Unidirectional coloring versus bicoloring



- Row (left) and column (right) coloring of an arrowhead matrix (center), both requiring the same number of colors as the matrix dimension (50 in this case).

Unidirectional coloring versus bicoloring



- Bicoloring of an arrowhead matrix, requiring 10 colors for the rows (left) and 10 colors for the columns (right).

Coloring and bicoloring on a rectangle matrix



Figure – Row coloring (left) and column coloring (right) of a rectangle matrix, requiring the same number of colors as the matrix dimensions (respectively 6 and 12 in this case).

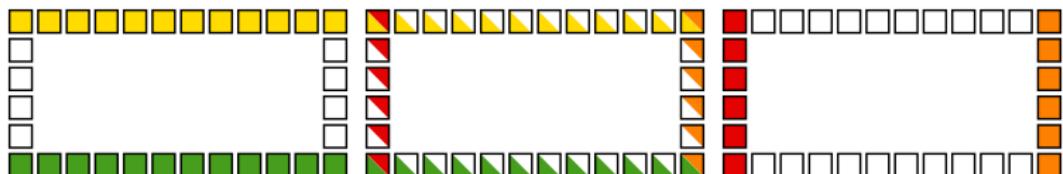


Figure – Bicoloring of a rectangle matrix, requiring only 2 colors for the rows (left) and 2 colors for the columns (right). In the central figure, each nonzero coefficient is colored using its row's color and its column's color.

- Gauss-Newton subproblem : $\min_{d \in \mathbb{R}^n} \|J(x_k)d + F(x_k)\|^2$ with $J \in \mathbb{R}^{m \times n}$.
- Column coloring + forward-mode AD is efficient when $m \gg n$.
- Dense rows (e.g., normalization constraints) make column coloring inefficient : entire row must be recovered.
- Row coloring + reverse-mode AD is inefficient if $m \gg n$ (many row colors).
- **Bicoloring** : recover sparse columns with forward-mode, few dense rows with reverse-mode → improved performance.

- Consider $\min_{x \in \mathbb{R}^n} f(x)$ subject to $c(x) = 0$
- $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $c : \mathbb{R}^n \rightarrow \mathbb{R}^m$, and the Jacobian $J_c(x) \in \mathbb{R}^{m \times n}$
- Row coloring + reverse-mode AD is efficient when $n \gg m$.
- Dense columns (variables affecting many constraints) make row coloring inefficient.
- Column coloring + forward-mode AD is an alternative but may be suboptimal.
- **Bicoloring**: recover sparse rows with reverse-mode, dense columns with forward-mode → better overall efficiency.

How to perform a bidirectional coloring ?

- Bicoloring and symmetric coloring share similarities.
- Bicoloring : Recover coefficients from rows or columns.
- Symmetric coloring : Recover coefficients from upper or lower triangle.
- Can we use star and acyclic symmetric colorings for bicoloring?

How to perform a bidirectional coloring?

- Bicoloring on a Jacobian J can be seen as a symmetric coloring on $H = \begin{bmatrix} 0 & J^T \\ J & 0 \end{bmatrix}$.
- We can easily derive both **direct** (star) and **substitution** (acyclic) bicoloring.

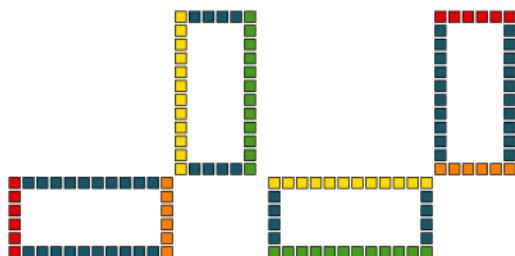


Figure – Symmetric coloring on H . Nonzeros are colored by the color of their columns on the left panel and by the color of their rows on the right panel.

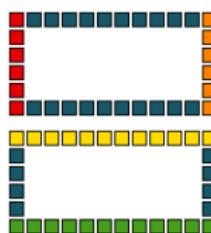


Figure – Bicoloring on J . Nonzeros are colored according to their column colors in the top panel and according to their row colors in the bottom panel.

Relation between neutral color and two-colored structures

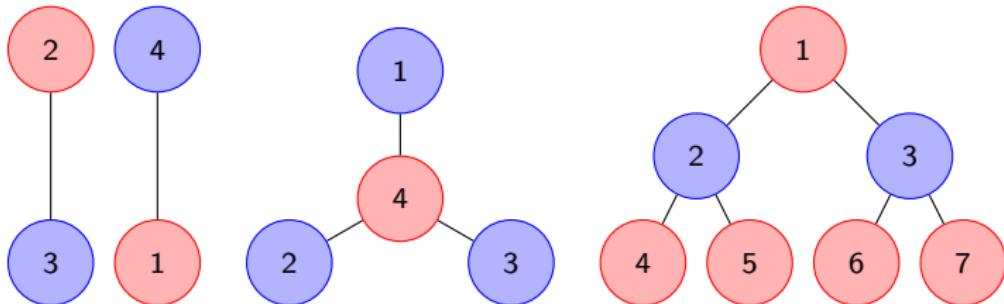


Figure – Variants of two-colored structures with trivial stars and trees (left), normal star (center) and normal tree (right).

- Diagonal entries take the color of their column, but are always zero under bicoloring.
- Normal trees require both colors for decompression.
- In normal stars, spoke colors are irrelevant for decompression.
- For trivial structures, the decompression color may be chosen arbitrarily from either vertex.

Content of SparseMatrixColorings.jl

SparseMatrixColorings.jl is a registered Julia package dedicated to coloring sparse Jacobians and Hessians.

```
pkg> add SparseMatrixColorings  
julia> using SparseMatrixColorings
```

SparseMatrixColorings.jl implements algorithms from our research and the following articles :

- *What Color Is Your Jacobian? Graph Coloring for Computing Derivatives*, Gebremedhin et al. (2005)
- *New Acyclic and Star Coloring Algorithms with Application to Computing Hessians*, Gebremedhin et al. (2007)
- *Efficient Computation of Sparse Hessians Using Coloring and Automatic Differentiation*, Gebremedhin et al. (2009)
- *ColPack : Software for graph coloring and related problems in scientific computing*, Gebremedhin et al. (2013)
- *Revisiting sparse matrix coloring and bicoloring*, Montoison et al. (2025)

Content of SparseMatrixColorings.jl

The three main functions to perform a coloring are `coloring`, `ColoringProblem` and `GreedyColoringAlgorithm`.

```
using SparseMatrixColorings, SparseArrays

S = sparse([
    1 1 1 1 1 1 1 1 1 1
    1 0 0 0 0 0 0 0 0 1
    1 0 0 0 0 0 0 0 0 1
    1 0 0 0 0 0 0 0 0 1
    1 1 1 1 1 1 1 1 1 1
])
problem = ColoringProblem(; structure=:nonsymmetric,
                           partition=:bidirectional)

order = RandomOrder()

algo = GreedyColoringAlgorithm(order;
                                decompression=:direct,
                                postprocessing=true)

result = coloring(S, problem, algo)
```

Content of SparseMatrixColorings.jl

Based on the result of `coloring`, you can easily recover a vector of integer colors with `row_colors`, `column_colors`, as well as the groups of colors with `row_groups` and `column_groups`.

```
julia> column_colors(result)
1
0
0
0
0
0
0
0
0
0
0
0
0
0
0
2
```

```
julia> column_groups(result)
[1]
[10]
```

Content of SparseMatrixColorings.jl

```
julia> row_colors(result)  
2  
0  
0  
0  
1
```

```
julia> row_groups(result)  
[5]  
[1]
```

```
julia> ncolors(result)  
4
```

Content of SparseMatrixColorings.jl

The functions `compress` and `decompress` efficiently store and retrieve compressed representations of colorings for sparse matrices.

```
A = sparse([
    1  2  3  4  5  6  7  8  9  10
    11 0  0  0  0  0  0  0  0  14
    12 0  0  0  0  0  0  0  0  15
    13 0  0  0  0  0  0  0  0  16
    17 18 19 20 21 22 23 24 25 26
])
```

```
Br, Bc = compress(A, result)
2×10 Matrix{Int64}:
 17  18  19  20  21  22  23  24  25  26
  1    2    3    4    5    6    7    8    9   10
```

```
5×2 Matrix{Int64}:
 1  10
 11 14
 12 15
 13 16
 17 26
```

Content of SparseMatrixColorings.jl

```
julia> C = decompress(Br, Bc, result)
5×10 SparseMatrixCSC{Int64, Int64} with 26 stored entries:
 1   2   3   4   5   6   7   8   9   10
11   .   .   .   .   .   .   .   .   14
12   .   .   .   .   .   .   .   .   15
13   .   .   .   .   .   .   .   .   16
17   18   19   20   21   22   23   24   25   26
```

```
julia> decompress!(A, 2*Br, 3*Bc, result)
5×10 SparseMatrixCSC{Int64, Int64} with 26 stored entries:
 6   8   12   16   20   24   28   32   36   60
66   .   .   .   .   .   .   .   .   84
72   .   .   .   .   .   .   .   .   90
78   .   .   .   .   .   .   .   .   96
102  72   76   80   84   88   92   96   100  156
```

What actually matters for JuMP and MOI?

- Jacobian coloring is unnecessary: expression trees already enable very efficient reverse-mode passes.
- But recent work on bicoloring introduces the idea of neutral colors in symmetric colorings and post-processing.
- These neutral colors become directly useful in MOI if we stop assuming a fully nonzero Hessian diagonal.

Ordering strategies matter

- MOI only supports the natural ordering of vertices.
- Many vertex orderings exist (random, largest first, smallest last, incidence degree, dynamic largest first, ...) and produce different colorings.
- We can precompute multiple colorings with different orderings as a preprocessing phase.
- Perfect elimination ordering is optimal for acyclic coloring on banded matrices or matrices with chordal-like sparsity.

Acyclic vs. star coloring

- MOI currently relies on acyclic coloring.
- Only the colors are kept; tree structures are discarded, limiting efficient preparation for decompression.
- Star coloring is cheaper to compute but yields more colors. We can alternate decompression and directional derivatives without storing all compressed Hessian columns.
- No need for `DataStructures.IntDisjointSets`: the forest structure in `SparseMatrixColorings.jl` already captures everything, and could replace the `DataStructures.jl` dependency in MOI.

- Neutral colors can be used in symmetric coloring for generic Hessian AD.
- Multiple-coloring preprocessing could improve robustness and reduce AD passes.
- Integration of SparseMatrixColorings.jl in MOI, potentially inside a new AD backend ?

Optimal coloring with JuMP / MOI

We implemented an optimal column / row coloring algorithm based on constraint programming in JuMP.

```
n = nb_vertices(bipartite_graph, Val(side))
model = Model(optimizer)

# one variable per vertex to color, removing some renumbering
# symmetries
@variable(model, 1 <= color[i=1:n] <= i, Int)

# one variable to count the number of distinct colors
@variable(model, ncolors, Int)
@constraint(model, [ncolors; color] in MOI.CountDistinct(n + 1))

# neighbors of the same vertex must have distinct colors
for i in vertices(bg, Val(other_side))
    neigh = neighbors(bg, Val(other_side), i)
    @constraint(model, color[neigh] in
        MOI.AllDifferent(length(neigh)))
end

# minimize the number of distinct colors
@objective(model, Min, ncolors)
optimize!(model)
```

Optimal coloring with JuMP / MOI

Still need to add a JuMP formulation for symmetric colorings.

```
using SparseMatrixColorings, JuMP, MathOptInterface, MiniZinc

coloring_problem = ColoringProblem(;  
    structure=:nonsymmetric, partition=:column)

algo = OptimalColoringAlgorithm(  
() -> MiniZinc.Optimizer{Float64}("highs");  
    silent=false, assert_solved=false)

coloring(J, coloring_problem, algo)
num_colors = ncolors(result)

import ORTools_jll
path_cp_sat = joinpath(ORTools_jll.artifact_dir, "share",
    "minizinc", "solvers", "cp-sat.msc")

algo = OptimalColoringAlgorithm(()-> MiniZinc.Optimizer{Float64}
    (path_cp_sat); silent=false, assert_solved=false)

coloring(J, coloring_problem, algo)
num_colors = ncolors(result)
```

Downloading SparseMatrixColorings.jl



<https://github.com/gdalle/SparseMatrixColorings.jl>