

# **MathOptInterface: a comprehensive overview**

Oscar Dowson

JuMP-dev 2025



# MathOptInterface.jl

The interface between modeling languages and the solvers

JuMP

Convex.jl

Optimization.jl

...

MathOptInterface.jl

HiGHS.jl

SCS.jl

Ipopt.jl

...

HiGHS

SCS

Ipopt

...



# The purpose of this talk

## Why is this talk needed:

- MathOptInterface.jl (MOI) is one of the largest packages in all of Julia
- It is the connection between JuMP and solvers
- It uses a novel abstraction
- It has a looooooot of stuff in it

We haven't publicly talked about it much

## By the end of this talk you will:

1. Understand the problem we are trying to solve and why we wrote MOI
2. Understand the MOI abstraction
3. Understand what a bridge is and why they are necessary
4. Have an overview of the components in MathOptInterface.jl


## By the end of this talk you will not:

- Be able to write a solver wrapper
- Know how to write code that uses MOI



# 2013-2019: MathProgBase.jl

MPB divided the world into three problem classes





# 2013-2019: Problems with MathProgBase

There were many. Here are four.

## Standard forms are not standard

- ECOS and SCS use different orderings for the exponential cone
- Gurobi does not support  $Ax$  in Interval
- CSDP does not support free variables

## Some solvers mix problem classes

- KNITRO supported nonlinear, but also SecondOrderCone and Complements
- Gurobi was linear quadratic, but now also supports nonlinear

## Extending the classes is hard

- No indicator constraints
- No complementarity constraints


## In-place problem modification was limited

- No support for deleting variables or constraints
- Support for modifying RHS but not LHS



# The 2017 JuMP-dev workshop

## Release JuMP as 1.0 or rewrite from scratch?





# The MathOptInterface standard form

MOI defines a very *regular* standard form

minimize:  $f_0(x)$

subject to:  $f_i(x) \in \mathcal{S}_i \quad \forall i \in \{1, \dots, m\}$

---

Function	Set
$x$ , Binary	VariableIndex(x) ZeroOne()
$x \geq 0$	VariableIndex(x) GreaterThan(0.0)
$2x + 1 == y$	ScalarAffineFunction(2x - y + 0) EqualTo(-1.0)
$Ax \geq b$	VectorAffineFunction(Ax - b) Nonnegatives()
$X \succeq 0$ , PSD	VectorOfVariables(vec(X)) PositiveSemidefiniteConeSquare()



# Solvers support a subset of functions and sets

## HiGHS

### Constraints

```
ScalarAffineFunction in {  
    EqualTo, GreaterThan, Interval, LessThan  
}
```

```
VariableIndex in {  
    EqualTo, GreaterThan, Interval, LessThan,  
    Integer, ZeroOne, Semicontinuous,  
    Semiinteger,  
}
```

### Objective

```
ScalarAffineFunction  
ScalarQuadraticFunction
```



# Solvers support a subset of functions and sets

## SCS

### Constraints

```
VectorAffineFunction in {  
    Nonnegatives, Zeros, SecondOrderCone,  
    ExponentialCone, DualExponentialCone,  
    PowerCone, DualPowerCone, ScaledPSDCone,  
    NormNuclearCone, ScaledComplexPSDCone,  
    ScaledLogDetConeTriangle,  
}
```

### Objective

```
ScalarAffineFunction  
ScalarQuadraticFunction
```



# Solvers support a subset of functions and sets

## Ipopt

### Constraints

```
VariableIndex in {  
    GreaterThan, LessThan, EqualTo, Interval,  
    Parameter,  
}
```

```
ScalarAffineFunction in {  
    GreaterThan, LessThan, EqualTo, Interval,  
}
```

```
ScalarQuadraticFunction in {  
    GreaterThan, LessThan, EqualTo, Interval,  
}
```

### Constraints continued...

```
ScalarNonlinearFunction in {  
    GreaterThan, LessThan, EqualTo, Interval,  
}
```

```
VectorOfVariables in VectorNonlinearOracle
```

### Objective

```
VariableIndex  
ScalarAffineFunction  
ScalarQuadraticFunction  
ScalarNonlinearFunction
```



# User and solver might speak different formulations

## This is a problem. Should users rewrite their model for every solver?

There are multiple ways to write the same constraint

Scalar or vector constraints

- $x \in \text{GreaterThan}(0.0)$
- $[x] \in \text{Nonnegatives}(1)$
- $Ax == b$
- $Ax - b \in \text{Zeros}()$

Affine transformations from one set to another

- $[t, x, y] \in \text{SecondOrderCone}(2)$
- $[t, t / 2, x, y] \in \text{RotatedSecondOrderCone}(3)$

More complicated transforms

- $[t, X] \in \text{LogDetConeSquare}()$
- A set of constraints using PSD, ExponentialCone, and LessThan sets



# Constraint bridges

## Map between equivalent formulations

Each bridge takes as input:

- a function-in-set constraint

and it may:

- add new constraints of a different type
- add new decision variables

Each bridge also implements:

- `MOI.delete(model::MOI.ModelLike, bridge)`
- `MOI.get(model::MOI.ModelLike, ::MOI.ConstraintPrimal, bridge)`
- `MOI.set(model::MOI.ModelLike, ::MOI.ConstraintPrimalStart, bridge, value)`
- `MOI.get(model::MOI.ModelLike, ::MOI.ConstraintDual, bridge)`
- `MOI.set(model::MOI.ModelLike, ::MOI.ConstraintDualStart, bridge, value)`



# Bridges let you model unique constraint types

## “An affine expression is integer”

```
using JuMP, HiGHS
model = Model(HiGHS.Optimizer)
@variable(model, x)
@constraint(
    model,
    2 * x + 1 in MOI.Integer(),
)
```

```
julia> print_active_bridges(model)
* Unsupported constraint:
| MOI.ScalarAffineFunction-in-MOI.Integer
| bridged by:
|   MOIB.Constraint.ScalarSlackBridge
| may introduce:
|   * Supported constraint:
|     MOI.ScalarAffineFunction-in-MOI.EqualTo
|   * Supported variable: MOI.Integer
```

```
julia> print(unsafe_backend(model))
Feasibility
Subject to:
VariableIndex-in-Integer
v[2] ∈ ℤ
ScalarAffineFunction-in-EqualTo
0.0 + 2.0 x - 1.0 v[2] == -1.0
```



# Hypergraphs and shortest hyperpaths


In the ideal case, the solver supports the constraint





# Hypergraphs and shortest hyperpaths


The solver may support the constraint via a single transformation





# Hypergraphs and shortest hyperpaths


A chain of bridges may be needed, that introduce more than one constraint





# Hypergraphs and shortest hyperpaths


There may be many equivalent paths through the graph





# Hypergraphs and shortest hyperpaths

In practice, the graph can be arbitrarily complicated





# Hypergraphs and shortest hyperpaths

## The optimal bridge is the shortest hyperpath

Let the graph  $G = (N, E)$ , where:

- $N$  is the set of nodes
- $E$  is the set of bridges

Let  $S$  = the set of nodes supported by the solver

Each edge  $e$  has:

- A source node  $s(e)$  in  $N$
- A set of target nodes  $T(e)$  subset  $N$
- A weight  $w(e)$

Most bridges choose  $w(e) = 1$ . Some bridges  
use  $w(e) = 10$


Minimum hyperpath is the set of bridges  $e$  that minimize:

```
function C(n)
    if n in S
        return 0
    end
    return minimum(
        w(e) + sum(C(m) for m in T(e))
        for e in E if s(e) == n
    )
end
```



# Hypergraphs and shortest hyperpaths

## The optimal bridge is the shortest hyperpath





# There are other types of bridges

a.k.a. it's even more complicated than it seems

Objective bridges reformulate objective functions

User writes


$$\min x^2$$

Bridges.Objective.ScalarSlackBridge

$$\min t: t \geq x^2$$

Bridges.Constraint.QuadToSOC

$$\min t: [t, 1/2, x] \text{ in RotatedSecondOrderCone}()$$





# There are other types of bridges


## a.k.a. it's even more complicated than it seems

Variable bridges reformulate variable domains

Seems simple. Take a **variable -in-set** and replace by **variable -in- different set** with an **affine substitution rule**.

Every **x** in the model must be replaced by the substitution rule, and every solution **y** must be inverted back to the original model.

**Keeping track of substitutions is THE most complicated code in all of MOI.**





jump-dev / MathOptInterface.jl

Type  to search

Code Issues Pull requests Discussions Actions Security Insights Settings

There a  
a.k.a. It's e

Variable bridges

Seems simple  
by variable -in-  
substitution rule

Every x in the  
substitution rule  
inverted back

Keeping track  
complicated

## Update bridge optimizers with variable bridges #816

Merged

blegat merged 2 commits into master from bl/bridge\_opt\_with\_var on Aug 6, 2019

Conversation 7

Commits 2

Checks 0

Files changed 8



blegat commented on Aug 6, 2019

Member

...

Extracted from #759



Update bridge optimizers with variable bridges

831efd5

blegat force-pushed the bl/bridge\_opt\_with\_var branch from 9cca68c to 831efd5 6 years ago

Compare



mlubin approved these changes on Aug 6, 2019

View reviewed changes

mlubin left a comment

Member

...

I don't have any substantive comments, so merge when ready.




Solver

y in  
 $(\mathbb{R}_+, \mathbb{R}_+)$



# Hypergraphs and shortest hyperpaths


Complexity is hidden from the user





# Hypergraph

In practice, the





# MathOptInterface.jl is a monolith

122,000 lines of code

/src	79,000 lines	/test	43,000 lines
/src/Bridges	19,800 lines	/test/Bridges	18,100 lines
/src/FileFormats	7,500 lines		
/src/Nonlinear	6,100 lines		
/src/Test	22,300 lines	/test/FileFormats	8,200 lines
/src/Utilities	14,800 lines	/test/Utilities	11,200 lines



# MOI.

## The top-level API

MOI uses multiple dispatch.

A lot.

Every solver is a new type.

Every function is a new type.

Every set is a new type.

There are a small number of public functions.

Testing takes a long time. Running the tests creates 6000 methods and 378,411 MethodInstances.

```
abstract type ModelLike end
abstract type AbstractOptimizer <: ModelLike end
abstract type AbstractFunction end
abstract type AbstractSet end
```

```
empty!(::ModelLike)
is_empty(::ModelLike)
```

```
add_variable(::ModelLike)::VariableIndex
```

```
add_constraint(
    ::ModelLike,
    ::AbstractFunction,
    ::AbstractSet,
)::ConstraintIndex
```

```
optimize!(::ModelLike)
```



# MOI.

## Functions

```
abstract type AbstractFunction <: MA.AbstractMutable end
abstract type AbstractScalarFunction <: AbstractFunction end
abstract type AbstractVectorFunction <: AbstractFunction end

struct VariableIndex <: AbstractScalarFunction
    value::Int64
end

struct ScalarAffineTerm{T}
    coefficient::T
    variable::VariableIndex
end

mutable struct ScalarAffineFunction{T} <: AbstractScalarFunction
    terms::Vector{ScalarAffineTerm{T}}
    constant::T
end
```



# MOI.

## Sets

```
abstract type AbstractSet end
abstract type AbstractScalarSet <: AbstractSet end
abstract type AbstractVectorSet <: AbstractSet end

struct EqualTo{T<:Number} <: AbstractScalarSet
    value::T
end

struct SecondOrderCone <: AbstractVectorSet
    dimension::Int
end

struct SOS1{T<:Real} <: AbstractVectorSet
    weights::Vector{T}
end

struct Indicator{A,S<:AbstractScalarSet} <: AbstractVectorSet
    set::S
end
```



# MOI.Bridges

## Benoît's Bridge Bellman Bedlam

This submodule contains:

- Three submodules for the different types of bridges: Bridges.Constraint, Bridges.Objective, and Bridges.Variable
- 76 bridges, each of which may cover many function-in-set combinations
- Bridges.LazyBridgeOptimizer, with Bellman-Ford and variable bridge substitutions
- Some other stuff, like tests for bridges

Users don't need to know the inner details

```
MOI.instantiate(  
    HiGHS.Optimizer;  
    with_bridge_type = Float64,  
)
```

```
MOI.Bridges.full_bridge_optimizer(  
    HiGHS.Optimizer(),  
    Float64,  
)
```



# MOI.Bridges

## The logic is horrific, so we try to test a lot

Testing is hard because of the number of possible permutations of supported variables, constraints, and objectives.

We test:

- every bridge in isolation
- a very wide range of unit tests
- every solver before every release, and sometimes with every PR

Sometimes this is still not enough. There are many edges cases.

Each bridge runs a common suite of tests

```
MOI.Bridges.runtests(  
    MOI.Bridges.Constraint.ZeroOneBridge,  
    """  
        variables: x  
        x in ZeroOne()  
        """,  
        """  
        variables: x  
        x in Integer()  
        1.0 * x in Interval(0.0, 1.0)  
        """,  
)
```

Test Summary:	Pass	Total	Time
Bridges.runtests	32	32	0.0s



# MOI.FileFormats

## Read and write models to disk

This submodule contains six submodules:

- CBF: the .cbf format
- LP: the .lp format
- MOF: the .mof.json format
- MPS: the .mps format
- NL: the .nl format
- SDPA: the .sdpa format

Each submodule implements

- A new Model <: MOI.ModelLike
- Base.read!(io::IO, ::Model)
- Base.write(io::IO, ::Model)

### Downside to this design

To read and write models from disk, JuMP does a variation of:

```
function JuMP.write_to_file(  
    model::Model, filename::String)  
    dest = MOI.FileFormats.Model(; filename)  
    MOI.copy_to(dest, model)  
    MOI.write_to_file(dest, filename)  
    return  
end
```

We first have to create a copy of the entire model. This is slower than writing from the existing model object



# MOI.FileFormats.MathOptFormat

<https://jump.dev/MathOptFormat/>

```
{  
    "version": {"major": 1, "minor": 4},  
    "variables": [{"name": "x"}],  
    "objective": {  
        "sense": "min",  
        "function": {  
            "type": "ScalarAffineFunction",  
            "terms": [{"coefficient": 2, "variable": "x"}],  
            "constant": 1  
        }  
    },  
    "constraints": [{  
        "function": {"type": "Variable", "name": "x"},  
        "set": {"type": "GreaterThan", "lower": 1}  
    }]  
}
```

*MathOptFormat supports every model you can write in MOI. And it has a schema.*




# MOI.Nonlinear

## A lot of complicated stuff

This submodule contains:

- Nonlinear.Model
  - A nonlinear modeling interface
- ReverseAD
  - A library for sparse reverse-mode automatic differentiation
- SymbolicAD
  - A symbolic differentiation library

See my JuMP-dev 2022 and 2023 talks.



Improving Nonlinear Programming Support in JuMP | Oscar Dowson | JuliaCon 2023

Oscar Dowson  
JuMP-dev 2023

Improving nonlinear programming support in



Improving Nonlinear Programming Support in JuMP | Oscar Dowson | JuliaCon 2022

Oscar Dowson  
JuMP-dev 2022



# MOI.Test

## Test-Driven-Development of solver wrappers

This submodule contains

- A suite of >500 solver-independent test functions
- Accessible via `MOI.Test.runtests`
- Ability to include/exclude tests, adjust tolerances, etc

### Downsides

- This (c|sh)ould have been a separate package
- Calling it `Test` was a mistake because it conflicts with `Base.Test`

Solvers get access to thousands of tests

```
using Test, HiGHS
import MathOptInterface as MOI
@testset "runtests" begin
    MOI.Test.runtests(
        MOI.instantiate(
            HiGHS.Optimizer,
            with_bridge_type = Float64,
        ),
        MOI.Test.Config(; atol = 1e-7),
    )
end
```

Test Summary:	Pass	Total	Time
runtests	3237	3237	4m23.8s



# MOI.Utilities

## A useful dumping ground for random stuff

This submodule contains a *looooot* of stuff

- Utilities for working with functions: comparing them, creating them, modifying them
- **Utilities.MockOptimizer**: a fake optimizer to mock tests in MOI and JuMP
- **Utilities.CachingOptimizer**: an abstraction across how solvers handle incremental modification
- **Utilities.Model** and **Utilities.UniversalFallback**: model objects that support everything except optimize!
- **Utilities.CleverDicts**: a dict that starts as a Base.Vector but switches to a Base.Dict on deletion
- **Utilities.GenericModel**: a modular system for defining matrix-based storage

Just sooooo much other stuff...



# MOI.Utilities.CachingOptimizer

Abstract over differences in incremental modification

**MOI.Utilities.CachingOptimizer:**

- maintains two copies of the model:
  - model\_cache
  - optimizer
- is in one of three states
  - NO\_OPTIMIZER
  - EMPTY\_OPTIMIZER
  - ATTACHED\_OPTIMIZER

MOI.Utilities.CachingOptimizer


model\_cache::MOI.ModelLike

optimizer::MOI.AbstractOptimizer



# MOI.Utilities.CachingOptimizer


## Flow for adding a constraint





# MOI.Utilities.CachingOptimizer


## Flow for deleting a variable





# MOI.Utilities.CachingOptimizer

## Flow for calling optimize!





# 2013-2019: Problems with MathProgBase

There were many. Here are four.

## Standard forms are not standard

- ECOS and SCS use different orderings for the exponential cone
- Gurobi does not support  $Ax$  in Interval
- CSDP does not support free variables

## Some solvers mix problem classes

- KNITRO supported nonlinear, but also SecondOrderCone and Complements
- Gurobi was linear quadratic, but now also supports nonlinear

## Extending the classes is hard

- No indicator constraints
- No complementarity constraints

## In-place problem modification was limited

- No support for deleting variables or constraints
- Support for modifying RHS but not LHS



# Problems with MathOptInterface

There are many. Here are four.

## Arbitrary Indices

Each variable and constraint has an associated value::Int64. These do not need to be ordered and do not need to be contiguous. Constraints of different types can have the same value. Make them ordered by creating and unique by variable/constraint.

## AbstractVectorFunction

Why didn't we just make it  
Base.Vector{<:AbstractScalarFunction}?

These problems are too breaking for us to ever consider changing. There will not be a MathOptInterface 2.0. We don't want Python 2->3.  
But if you're looking to re-implement MOI in a different language... don't copy it blindly. Come talk to me.

## The 0.5 in ScalarQuadraticFunction

It's so hard to remember whether to \* or / the 0.5. Just make the function a list of terms. Not a Q matrix.

## Variable bridges (and constrained variables)

We haven't talked about add\_constrained\_variable. Variable bridges are hard because variable sets can overlap: @variable(model, x >= 0, Int)  
Is the domain of x Real, Nonnegatives, or Integer?



# The purpose of this talk

## Why is this talk needed:

- MathOptInterface.jl (MOI) is one of the largest packages in all of Julia
- It is the connection between JuMP and solvers
- It uses a novel abstraction
- It has a looooooot of stuff in it

We haven't publicly talked about it much

## By the end of this talk you will:

1. Understand the problem we are trying to solve and why we wrote MOI
2. Understand the MOI abstraction
3. Understand what a bridge is and why they are necessary
4. Have an overview of the components in MathOptInterface.jl

## By the end of this talk you will not:

- Be able to write a solver wrapper
- Know how to write code that uses MOI