



The life and times of SDDP.jl

<https://sddp.dev>

Oscar Dowson

JuMP-dev 2025



The purpose of this talk

Why is this talk needed:

- SDDP.jl is a somewhat popular (if niche) JuMP extension
- It has been in continuous development for 10 years
- It has a lot of ideas that might be useful for other extensions

I've never given a talk about it at JuMP-dev (or on YouTube)

By the end of this talk you will:

1. Have some knowledge of the New Zealand dairy and electricity industries
2. Understand the policy graph decomposition for modeling sequential decision problems
3. Know how SDDP.jl implements a JuMP extension, uses multiple dispatch, and supports multithreading


By the end of this talk you will not:

- Know the details of the SDDP algorithm



The New Zealand dairy industry


It's a stochastic optimal control problem





The New Zealand dairy industry


It's a stochastic optimal control problem





The New Zealand dairy industry


It's a stochastic optimal control problem





The New Zealand dairy industry

It's a stochastic optimal control problem





The New Zealand energy system

It's a stochastic optimal control problem



State variables

- Volume of water in each reservoir

Control variables

- Hydro generation
- Thermal generation

Random variables

- Inflow

Constraints

- water in = water out
- supply = demand

Objective function

- Minimize cost



Oscar Dowson

JuliaCon 2017



0:01 / 7:57 • Welcome! ▶



THE UNIVERSITY OF
AUCKLAND
Te Whare Wānanga o Tāmaki Makaurau
NEW ZEALAND

Cows, Lakes, and a JuMP extension for multi-stage stochastic programming

Oscar Dowson
The University of Auckland
odow003@aucklanduni.ac.nz



Cows, Lakes, and a JuMP Extension
for Multi-stage Stochastic Optimization



Joaquim Garcia

JuliaCon 2017



SDDP – Power system operation modeling

► Physical parameters

- Hydro (detailed topology (cascades), hydro production, reservoirs modeling, operative constraints etc.)
- Thermal (efficiency curves, combined cycle plants, multiple fuel plants, fuel availability constraints, GHG emission factors, unit commitment decisions etc.)
- Renewables (Wind, biomass, solar etc. represented scenarios)
- Transmission Network (Linearized power flow model with quadratic losses, security constraints etc.)

► Stochastic parameters

- Hydro inflows and renewable generation - Multivariate stochastic model
- Uncertainty on fuel costs - Markov chains (hybrid SDDP/SDP model)
- Wholesale energy market prices - Markov chains
- Generation & transmission equipment outages - Monte Carlo



Stochastic Optimization Models on Power Systems



6:42 / 35:31 • Help us add time stamps or captions to this video! See the description for details.





JADE: Just Another DOASA Environment

Contributions from the Electric Power Optimization Centre

1991: SDDP introduced by Pereira and Pinto

2008: Philpott & Guan wrote the first AMPL version of NZ model called DOASA

2012-16: Philpott and de Matos wrote a C++ version of DOASA

2016-18: I wrote SDDP.jl (for cows), Lea Kapelevich wrote a NZ energy model called JADE

2022: JADE was adopted by New Zealand electricity authority

The screenshot shows a web browser window with the URL 'emi.ea.govt.nz' in the address bar. The page header includes the EMI logo, 'ELECTRICITY AUTHORITY TE MANA HIKO', and navigation links for 'HOME', 'RETAIL', 'WHOLESALE', 'FORWARD MARKETS', 'ENVIRONMENT', and 'MY DASHBOARDS'. Below the header, a breadcrumb trail shows 'wholesale category > Tools > JADE'. The main content area is titled 'JADE overview' and describes JADE as a modelling package for the New Zealand electricity generation sector, mentioning its multistage stochastic optimization and hydrological aspects. It also highlights the presence of uncertainty and variability.

JADE overview

JADE is a modelling package that implements a multistage stochastic optimization representing the New Zealand electricity generation sector, with a rich treatment of the hydrological aspects of the sector. Key outputs of the model include a water value surface for each stage or week of the modelled time horizon, typically a year, and corresponding marginal water values for each reservoir represented in the model.

One of the difficulties with planning and operational decision making in a hydro-dominated electricity system such as New Zealand's is the uncertainty and variability associated with inflows into hydro storage reservoirs. JADE is an ideal tool to aid decision making in the presence of such uncertainty.

Some high-level characteristics of JADE:

- The [EPOC team](#) at Auckland University created and maintain the JADE modelling package. Significant contributions over the years have come from A Philpott, G Pritchard, A Downward, O Dowson, and L Kapelevich.
- JADE supersedes [DOASA](#), another EPOC model that the Authority has used for several years.
- JADE is formulated using the [JuMP](#) package, an algebraic modelling language for mathematical optimization written in the [Julia](#) programming language.
- At the heart of JADE is the Julia package for stochastic dual dynamic programming by Oscar Dowson, [SDDP.jl](#).
- JADE can be solved with open-source solvers, although a commercial solver requiring a paid license, e.g. Gurobi or Cplex, is recommended for large-scale models.
- JADE is open source and available from [GitHub](#).



Modelling


The ingredients

To model a policy graph we need

- A description of the graph
- A subproblem for each node

Each subproblem needs

1. Incoming and outgoing state variables
2. Control variables
3. Constraints
4. An objective function
5. A random variable





Modelling

An example

```
model = SDDP.LinearPolicyGraph(;  
    stages = 5, lower_bound = 0,  
) do sp::JuMP.Model, t::Int  
    @variable(sp, 0 <= x_reservoir <= 10, SDDP.State, initial_value = 5)  
    @variable(sp, u_thermal >= 0)  
    @variable(sp, u_hydro >= 0)  
    @constraint(sp, c_water, x_reservoir.out - x_reservoir.in + u_hydro <= 0)  
    @constraint(sp, c_demand, u_thermal + u_hydro == 1)  
    SDDP.@stageobjective(sp, t * u_thermal)  
    Ω, P = [0, 1, 2], [0.3, 0.5, 0.2]  
    SDDP.parameterize(sp, Ω, P) do ω  
        set_normalized_rhs(c_water, ω)  
    end  
end  
SDDP.train(model)  
simulations = SDDP.simulate(model, 100)
```






Modelling

A description of the graph

```
model = SDDP.LinearPolicyGraph(;  
    stages = 5, lower_bound = 0,  
) do sp::JuMP.Model, t::Int  
    @variable(sp, 0 <= x_reservoir <= 10, SDDP.State, initial_value = 5)  
    @variable(sp, u_thermal >= 0)  
    @variable(sp, u_hydro >= 0)  
    @constraint(sp, c_water, x_reservoir.out - x_reservoir.in + u_hydro <= 0)  
    @constraint(sp, c_demand, u_thermal + u_hydro == 1)  
    SDDP.@stageobjective(sp, t * u_thermal)  
    Ω, P = [0, 1, 2], [0.3, 0.5, 0.2]  
    SDDP.parameterize(sp, Ω, P) do ω  
        set_normalized_rhs(c_water, ω)  
    end  
end  
SDDP.train(model)  
simulations = SDDP.simulate(model, 100)
```





Modelling

State variables

```
model = SDDP.LinearPolicyGraph(;  
    stages = 5, lower_bound = 0,  
) do sp::JuMP.Model, t::Int  
    @variable(sp, 0 <= x_reservoir <= 10, SDDP.State, initial_value = 5)  
    @variable(sp, u_thermal >= 0)  
    @variable(sp, u_hydro >= 0)  
    @constraint(sp, c_water, x_reservoir.out - x_reservoir.in + u_hydro <= 0)  
    @constraint(sp, c_demand, u_thermal + u_hydro == 1)  
    SDDP.@stageobjective(sp, t * u_thermal)  
    Ω, P = [0, 1, 2], [0.3, 0.5, 0.2]  
    SDDP.parameterize(sp, Ω, P) do ω  
        set_normalized_rhs(c_water, ω)  
    end  
end  
SDDP.train(model)  
simulations = SDDP.simulate(model, 100)
```





Modelling

Control variables

```
model = SDDP.LinearPolicyGraph(;  
    stages = 5, lower_bound = 0,  
) do sp::JuMP.Model, t::Int  
    @variable(sp, 0 <= x_reservoir <= 10, SDDP.State, initial_value = 5)  
    @variable(sp, u_thermal >= 0)  
    @variable(sp, u_hydro >= 0)  
    @constraint(sp, c_water, x_reservoir.out - x_reservoir.in + u_hydro <= 0)  
    @constraint(sp, c_demand, u_thermal + u_hydro == 1)  
    SDDP.@stageobjective(sp, t * u_thermal)  
    Ω, P = [0, 1, 2], [0.3, 0.5, 0.2]  
    SDDP.parameterize(sp, Ω, P) do ω  
        set_normalized_rhs(c_water, ω)  
    end  
end  
SDDP.train(model)  
simulations = SDDP.simulate(model, 100)
```





Modelling Constraints

```
model = SDDP.LinearPolicyGraph(;  
    stages = 5, lower_bound = 0,  
) do sp::JuMP.Model, t::Int  
    @variable(sp, 0 <= x_reservoir <= 10, SDDP.State, initial_value = 5)  
    @variable(sp, u_thermal >= 0)  
    @variable(sp, u_hydro >= 0)  
    @constraint(sp, c_water, x_reservoir.out - x_reservoir.in + u_hydro <= 0)  
    @constraint(sp, c_demand, u_thermal + u_hydro == 1)  
    SDDP.@stageobjective(sp, t * u_thermal)  
    Ω, P = [0, 1, 2], [0.3, 0.5, 0.2]  
    SDDP.parameterize(sp, Ω, P) do ω  
        set_normalized_rhs(c_water, ω)  
    end  
end  
SDDP.train(model)  
simulations = SDDP.simulate(model, 100)
```





Modelling

Objective function

```
model = SDDP.LinearPolicyGraph(;  
    stages = 5, lower_bound = 0,  
) do sp::JuMP.Model, t::Int  
    @variable(sp, 0 <= x_reservoir <= 10, SDDP.State, initial_value = 5)  
    @variable(sp, u_thermal >= 0)  
    @variable(sp, u_hydro >= 0)  
    @constraint(sp, c_water, x_reservoir.out - x_reservoir.in + u_hydro <= 0)  
    @constraint(sp, c_demand, u_thermal + u_hydro == 1)  
    SDDP.@stageobjective(sp, t * u_thermal)  
    Ω, P = [0, 1, 2], [0.3, 0.5, 0.2]  
    SDDP.parameterize(sp, Ω, P) do ω  
        set_normalized_rhs(c_water, ω)  
    end  
end  
SDDP.train(model)  
simulations = SDDP.simulate(model, 100)
```





Modelling

Random variables

```
model = SDDP.LinearPolicyGraph(;  
    stages = 5, lower_bound = 0,  
) do sp::JuMP.Model, t::Int  
    @variable(sp, 0 <= x_reservoir <= 10, SDDP.State, initial_value = 5)  
    @variable(sp, u_thermal >= 0)  
    @variable(sp, u_hydro >= 0)  
    @constraint(sp, c_water, x_reservoir.out - x_reservoir.in + u_hydro <= 0)  
    @constraint(sp, c_demand, u_thermal + u_hydro == 1)  
    SDDP.@stageobjective(sp, t * u_thermal)  
  
    Ω, P = [0, 1, 2], [0.3, 0.5, 0.2]  
    SDDP.parameterize(sp, Ω, P) do ω  
        set_normalized_rhs(c_water, ω)  
    end  
end  
SDDP.train(model)  
simulations = SDDP.simulate(model, 100)
```





Modelling

All of it together

```
model = SDDP.LinearPolicyGraph(;  
    stages = 5, lower_bound = 0,  
) do sp::JuMP.Model, t::Int  
    @variable(sp, 0 <= x_reservoir <= 10, SDDP.State, initial_value = 5)  
    @variable(sp, u_thermal >= 0)  
    @variable(sp, u_hydro >= 0)  
    @constraint(sp, c_water, x_reservoir.out - x_reservoir.in + u_hydro <= 0)  
    @constraint(sp, c_demand, u_thermal + u_hydro == 1)  
    SDDP.@stageobjective(sp, t * u_thermal)  
    Ω, P = [0, 1, 2], [0.3, 0.5, 0.2]  
    SDDP.parameterize(sp, Ω, P) do ω  
        set_normalized_rhs(c_water, ω)  
    end  
end  
SDDP.train(model)  
simulations = SDDP.simulate(model, 100)
```






Solution algorithm

SDDP a.k.a each node is Benders

```
julia> SDDP.train(model; time_limit = 240)
```

```
-----  
SDDP.jl (c) Oscar Dowson and contributors, 2017-25  
-----  
problem  
nodes      : 108  
state variables : 5  
scenarios   : Inf  
options  
solver      : SDDP.Threaded()  
risk measure : SDDP.Expectation()  
-----  
iteration simulation bound      time (s)    solves  pid  
-----  
1 -1.415148e+04 1.000000e+06 3.980785e+00 33013  3  
2 -1.893675e+04 1.000000e+06 4.179685e+00 40286  1  
3 -1.033561e+04 1.000000e+06 4.842130e+00 66687  1  
† 4 -2.486776e+04 7.881972e+05 7.643921e+00 175091 1  
† 5 -2.831493e+05 7.588107e+05 1.909121e+01 514620 4  
6 -3.021992e+05 7.588107e+05 2.195494e+01 554152 2  
7 -1.556580e+05 7.308679e+05 2.195530e+01 554152 3  
8 -1.025802e+05 7.308679e+05 2.195624e+01 554155 1  
† 21 -6.585407e+04 2.692613e+05 6.075782e+01 1568680 3  
83 3.115513e+05 1.272015e+05 2.112643e+02 4539440
```



We modify and solve 4.5e6 LPs in 200 seconds.



Multithreading

Non-deterministic concurrent

Each thread iterates independently

There is a lock at each node

Works great if # nodes >> # threads

Problems

Gurobi environments are not thread-safe.

Need a separate license for each node (not each thread)

Limited tooling to detect race conditions

Trivial to implement

```
# Serial
while iteration(model, options)
end

# Parallel
interrupt = Threads.Atomic{Bool}(false)
@sync for _ in 1:Threads.threads()
    Threads.@spawn try
        while !interrupt[] && iteration(model, options)
            end
    finally
        interrupt[] = true
    end
end
```



Multiple dispatch

We use it. Perhaps too much

There are many “plug-ins” in SDDP.jl

| Type | <i>Controls...</i> |
|------------------------|---|
| AbstractRiskMeasure | How random variables are aggregated into a scalar |
| AbstractDualityHandler | How we compute the reduced of a fixed variable |
| AbstractSamplingScheme | How we sample trajectories in the graph |
| AbstractForwardPass | The forward pass |
| AbstractBackwardPass | The backward pass |



Multiple dispatch

SDDP.AbstractRiskMeasure

```
struct Expectation <: SDDP.AbstractRiskMeasure
end

function SDDP.adjust_probability(
    measure::Expectation, q, p, ω, X, is_min,
)
    q .= p
    return 0.0
end

SDDP.train(
    model;
    risk_measure = SDDP.Expectation(),
)
```

```
struct Entropic <: SDDP.AbstractRiskMeasure
    γ::Float64
end

function SDDP.adjust_probability(
    measure::Entropic, q, p, ω, X, is_min,
)
    γ = is_min ? measure.γ : -measure.γ
    y = p .* exp.(big.(γ .* X))
    q .= y / sum(y)
    return -q' * log.(q ./ p) / γ
end

SDDP.train(
    model;
    risk_measure = SDDP.Entropic(10.0),
)
```



Multiple dispatch

SDDP.AbstractDualityHandler

```
struct ContinuousConicDuality <:  
    SDDP.AbstractDualityHandler  
end  
function SDDP.get_dual_solution(  
    node::SDDP.Node, ::ContinuousConicDuality,  
)  
    undo =  
        relax_integrality(node.subproblem)  
    JuMP.optimize!(node.subproblem)  
    ret = Dict(  
        name => JuMP.dual(JuMP.FixRef(x.in))  
        for (name, x) in node.states  
    )  
    undo()  
    return ret  
end  
  
duality_handler = SDDP.ContinuousConicDuality()  
SDDP.train(model; duality_handler)
```

```
struct FixedDiscreteDuality <:  
    SDDP.AbstractDualityHandler  
end  
function SDDP.get_dual_solution(  
    node::SDDP.Node, ::FixedDiscreteDuality,  
)  
    undo =  
        fix_discrete_variables(node.subproblem)  
    JuMP.optimize!(node.subproblem)  
    ret = Dict(  
        name => JuMP.dual(JuMP.FixRef(x.in))  
        for (name, x) in node.states  
    )  
    undo()  
    return ret  
end  
  
duality_handler = SDDP.FixedDiscreteDuality()  
SDDP.train(model; duality_handler)
```



Takeaways

If you remember nothing else, go to <https://sddp.dev>

For modelers

- Sequential decision making under uncertainty is pervasive
- One approach is to model them as a policy graph
- We have general purpose software for solving policy graphs

For JuMP developers

- JuMP extensions allow custom syntax for users
- Writing multithreaded algorithms is “easy”
- Multiple dispatch makes it trivial to provide plugins that change the algorithm

restricted modeling + high quality general purpose software

>>>

generic modeling that requires custom algorithms